

Generative Adversarial Imitation Learning

Jonathan Ho, Stefano Ermon - NIPS 2016

Stanford University


Presenter: Jinsung Yoon

Problem: Imitation Learning

- Learning to perform a task from **demonstrations (Expert trajectories)** **without interaction** with the expert or access to reinforcement signal
- Applications: When the **rewards are hard to define**



O_t



$\pi_{\theta}(u_t | O_t)$



u_t

Examples:

- **Autonomous vehicles:**

<https://www.youtube.com/watch?v=cFtnflNe5fM>

- **Medicine**

Problem: Imitation Learning

Two main approaches

- **Behavioral cloning (BC)**


- A supervised learning problem that maps state/action pairs to policy.
(State: feature, action: label)
- Requires a large number of expert trajectories (**high sample complexity**) - due to **compounding error** caused by covariate shift
- Copies unnecessary actions as well

- **Inverse Reinforcement Learning (IRL)**

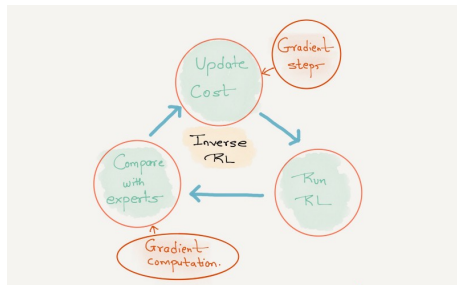
- Learns the **reward function** from expert trajectories that **prioritizes** entire trajectories over others, then derives the optimal policy
- Expensive to run (**Inner loop has RL**)
- Indirectly learns optimal policy from the reward function (**Using RL**)

Problem: Imitation Learning

- Behavioral cloning (BC)



- Inverse Reinforcement Learning (IRL)



Generative Adversarial Imitation Learning (GAIL)

- **Directly extracting a policy** from data as if it were obtained by RL following IRL.
- **Bypassing** any intermediate IRL step
- Draws an analogy between **imitation learning** and **generative adversarial networks (GAN)**
- Derive a model-free imitation learning algorithm with **significant performance improvement** with low sample and computational complexity.

Background: Preliminaries on RL

- \mathcal{S} : Finite state space
- \mathcal{A} : Finite action space
- Π : the set of all stochastic policies. Take action $\in \mathcal{A}$ given state $\in \mathcal{S}$
- $P(s'|s, a)$: Model dynamics
- $\pi \in \Pi$: A policy
- γ -discounted infinite horizon setting

$$\mathbb{E}_{\pi}[c(s, a)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t c(s_t, a_t)\right]$$

- where $s_0 \sim p_0, a_t \sim \pi(\cdot|s_t), s_{t+1} \sim P(\cdot|s_t, a_t)$
- π_E : Expert policy

Inverse Reinforcement Learning (IRL)

- **Maximum causal entropy IRL**

$$\max_{c \in \mathcal{C}} \left[\min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right] - \mathbb{E}_{\pi_E}[c(s, a)]$$

- where $H(\pi) = [E]_{\pi}[-\log \pi(a|s)]$
- Try to find a cost function $c \in \mathcal{C}$ that assigns **low cost** to the expert policy π_E and **high cost** to other policies (π)
- Using **RL procedure**, we can find the expert policy based on the cost c

$$RL(c) = \arg \min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)]$$

- Inner loop has **RL**; thus, slow

Proposed Framework

- Use the **largest possible set of cost functions**
 $\mathcal{C} = \mathbb{R}^{\mathcal{S} \times \mathcal{A}} = \{c : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}\}$
- Use Gaussian processes or **Neural networks** to find the best cost function c among large cost function class \mathcal{C}
- Avoid overfitting, we use a "**convex**" **cost function regularizer**
 $\psi : \mathbb{R}^{\mathcal{S} \times \mathcal{A}} \rightarrow \mathbb{R} \cup \infty$
- With ψ , IRL procedure can be written as

$$\text{IRL}_{\psi}(\pi_E) = \arg \max_{c \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}} -\psi(c) + \left[\min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right] - \mathbb{E}_{\pi_E}[c(s, a)]$$

- Let $\tilde{c} \in \text{IRL}_{\psi}(\pi_E)$, we are interested in $\pi = \text{RL}(\tilde{c})$

Occupancy Measure

- For a policy $\pi \in \Pi$, define its occupancy measure $\rho_\pi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ as

$$\rho_\pi(s, a) = \pi(a|s) \sum_{t=0}^{\infty} \gamma^t P(s_t = s | \pi)$$

- In words, occupancy measure is the **distribution** of state-action pairs with policy π

$$\mathbb{E}_\pi[c(s, a)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t c(s_t, a_t)\right] = \sum_{s, a} \rho_\pi(s, a) c(s, a)$$

- The set of valid occupancy measures $\mathcal{D} = \{\rho_\pi : \pi \in \Pi\}$ can be written as

$$\mathcal{D} = \left\{ \rho : \rho \geq 0 \& \sum_a \rho(s, a) = p_0(s) + \gamma \sum_{s', a} P(s|s', a) \rho(s', a) \forall s \in \mathcal{S} \right\}$$

- Note that there is **1-1 correspondence** between Π and \mathcal{D}

Occupancy Measure

- π_ρ to denote the **unique policy for an occupancy measure** ρ
- **Convex conjugate:** for a function $f : \mathbb{R}^{S \times \mathcal{A}} \rightarrow \mathbb{R} \cup \infty$, its convex conjugate $f^* : \mathbb{R}^{S \times \mathcal{A}} \rightarrow \mathbb{R} \cup \infty$ is

$$f^*(x) = \sup_{y \in \mathbb{R}^{S \times \mathcal{A}}} x^T y - f(y)$$

- Then, $\text{RL}(\tilde{c})$ can be written as

$$\text{RL} \odot \text{IRL}_\psi(\pi_E) = \arg \min_{\pi \in \Pi} -H(\pi) + \psi^*(\rho_\pi - \rho_{\pi_E})$$

- Proof:

$$\begin{aligned} & \arg \min_{\pi \in \Pi} -H(\pi) + \psi^*(\rho_\pi - \rho_{\pi_E}) \\ &= \arg \min_{\pi \in \Pi} \max_c -H(\pi) - \psi(c) + \sum_{s,a} \rho(s,a)c(s,a) - \sum_{s,a} \rho_{\pi_E}(s,a)c(s,a) \end{aligned}$$

Occupancy Measure

- The above theorem said that ψ -regularized IRL, implicitly seeks a **policy whose occupancy measure is close to the expert's** as measured by the convex function ψ^*
- It shows that various settings of ψ lead to various imitation learning algorithms that **directly solve the optimization problem**. (without IRL and RL iteration)
- Therefore we can deduce the following things
 - IRL is a dual of an **occupancy measure matching problem**
 - The induced optimal policy is the **primal optimum**.
- Therefore, the traditional IRL definition (**finding a cost function that the expert policy is uniquely optimal**) changes to (**finding a policy that matches the expert's occupancy measure**)

Generative Adversarial Imitation Learning (GAIL)

- Proposed ψ

$$\psi_{GA}(c) = \begin{cases} \mathbb{E}_{\pi_E}[g(c(s, a))], & \text{if } c < 0. \\ \infty, & \text{otherwise.} \end{cases}$$

where

$$g(x) = \begin{cases} -x - \log(1 - e^x), & \text{if } x < 0. \\ \infty, & \text{otherwise.} \end{cases}$$

- **Low penalty** when x is far from 0. **High penalty** when x is close to 0.

- The regularization function ψ_{GA} is **motivated by the following fact.**

$$\psi_{GA}^*(\rho_\pi - \rho_{\pi_E}) = \max_{D \in (0,1)^{\mathcal{S} \times \mathcal{A}}} \mathbb{E}_\pi[\log(D(s, a))] + \mathbb{E}_{\pi_E}[\log(1 - D(s, a))]$$

- where the **maximum ranges over discriminative classifiers** is $D : \mathcal{S} \times \mathcal{A} \rightarrow (0, 1)$
- The above equation is the **optimal negative log loss** of the binary classification problem of distinguishing between state-action pairs of π and π_E .
- Therefore, the optimal loss is the **Jensen-Shannon divergence**

$$D_{JS}(\rho_\pi, \rho_{\pi_E}) = D_{KL}(\rho_\pi \| (\rho_\pi + \rho_E)/2) + D_{KL}(\rho_{\pi_E} \| (\rho_\pi + \rho_{\pi_E})/2)$$

Proposed Optimization Problem

$$\min_{\pi} \psi_{\text{GA}}^*(\rho_{\pi} - \rho_{\pi_E}) - \lambda H(\pi) = D_{\text{JS}}(\rho_{\pi}, \rho_{\pi_E}) - \lambda H(\pi)$$

- It finds a policy whose **occupancy measure minimizes Jensen-Shannon divergence to the expert's.**

Proposed Algorithm

$$\min_{\pi} \max_{D \in (0,1)^{\mathcal{S} \times \mathcal{A}}} \mathbb{E}_{\pi}[\log(D(s, a))] + \mathbb{E}_{\pi_E}[\log(1 - D(s, a))] - \lambda H(\pi)$$

- We find a saddle point (π, D)

$$\min_{\pi} \max_{D \in (0,1)^{\mathcal{S} \times \mathcal{A}}} \mathbb{E}_{\pi}[\log(D(s, a))] + \mathbb{E}_{\pi_E}[\log(1 - D(s, a))] - \lambda H(\pi)$$

- Initialize the **policy** π_{θ} , and a **discriminator** $D_w : \mathcal{S} \times \mathcal{A} \rightarrow (0, 1)$
- **Alternatively update** θ and w
 - **Adam for gradient step on** w to increase $\mathbb{E}_{\pi}[\log(D(s, a))] + \mathbb{E}_{\pi_E}[\log(1 - D(s, a))]$
 - **TPRO step on** θ to decrease $\mathbb{E}_{\pi}[\log(D(s, a))] + \mathbb{E}_{\pi_E}[\log(1 - D(s, a))] - \lambda H(\pi)$
- Discriminator network is a **local cost function** providing learning signal to the policy.
- Taking a policy step that **decreases expected cost** w.r.t $c(s, a) = \log D(s, a)$

Try to find the policy $\pi \in \Pi$ that minimizes the cost function $c(s, a)$

$$L(\theta) = \mathbb{E}_{(s,a) \sim p}[c(s, a)]$$

Therefore, its gradient is

$$\begin{aligned}\nabla_{\theta} L(\theta) &= \nabla_{\theta} \int_{s,a} c(s, a) p(s, a) \\ &= \nabla_{\theta} \int_{s,a} c(s, a) \pi_{\theta}(a|s) p(s) \\ &= \int_{s,a} c(s, a) \nabla_{\theta} \pi_{\theta}(a|s) p(s) \\ &= \int_{s,a} c(s, a) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)} \pi_{\theta}(a|s) p(s) \\ &= \int_{s,a} c(s, a) \nabla_{\theta} \log(\pi_{\theta}(a|s)) \pi_{\theta}(a|s) p(s) \\ &= \mathbb{E}_{(s,a) \sim p}[c(s, a) \nabla_{\theta} \log(\pi_{\theta}(a|s))]\end{aligned}$$

Algorithm 1 Generative adversarial imitation learning

- 1: **Input:** Expert trajectories $\tau_E \sim \pi_E$, initial policy and discriminator parameters θ_0, w_0
- 2: **for** $i = 0, 1, 2, \dots$ **do**
- 3: Sample trajectories $\tau_i \sim \pi_{\theta_i}$
- 4: Update the discriminator parameters from w_i to w_{i+1} with the gradient

$$\hat{\mathbb{E}}_{\tau_i} [\nabla_w \log(D_w(s, a))] + \hat{\mathbb{E}}_{\tau_E} [\nabla_w \log(1 - D_w(s, a))] \quad (17)$$

- 5: Take a policy step from θ_i to θ_{i+1} , using the TRPO rule with cost function $\log(D_{w_{i+1}}(s, a))$. Specifically, take a KL-constrained natural gradient step with

$$\begin{aligned} & \hat{\mathbb{E}}_{\tau_i} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q(s, a)] - \lambda \nabla_{\theta} H(\pi_{\theta}), \\ & \text{where } Q(\bar{s}, \bar{a}) = \hat{\mathbb{E}}_{\tau_i} [\log(D_{w_{i+1}}(s, a)) \mid s_0 = \bar{s}, a_0 = \bar{a}] \end{aligned} \quad (18)$$

- 6: **end for**
-

Experiments

Experimental Settings (Run on OpenAI Gym)

- **Low-dimensional** control tasks: (e.g. Cartpole, Acrobot)
- **High-dimensional** tasks: (e.g. 3D humanoid locomotion)

Procedures

- Generate expert behavior for these tasks by **running TRPO** on the **true cost functions** to create expert policies.
- Run GAIL and other benchmarks on the **generated expert policies**.
- Evaluate imitation performance w.r.t **sample complexity of expert data**.

Benchmarks

- Behavior Cloning
- Feature expectation matching (FEM): with **linear** cost function
- Game-theoretic apprenticeship learning (GTAL): with **convex** cost function

Results

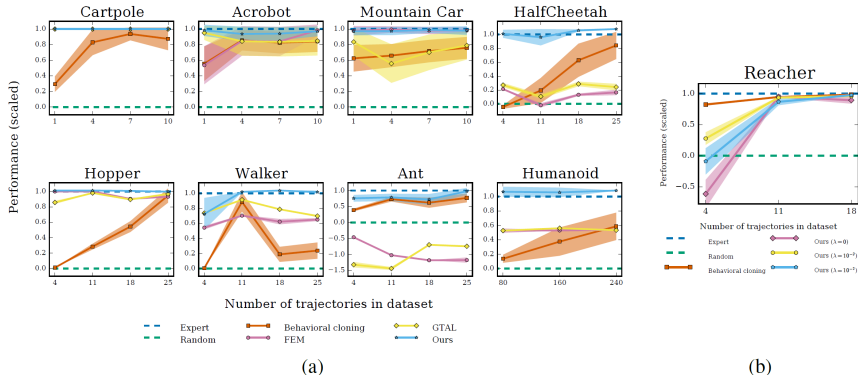


Figure 1: (a) Performance of learned policies. The y -axis is negative cost, scaled so that the expert achieves 1 and a random policy achieves 0. (b) Causal entropy regularization λ on Reacher.

Table 3: Learned policy performance

Task	Dataset size	Behavioral cloning	FEM	GTAL	Ours
Cartpole	1	72.02 \pm 35.82	200.00 \pm 0.00	200.00 \pm 0.00	200.00 \pm 0.00
	4	169.18 \pm 59.81	200.00 \pm 0.00	200.00 \pm 0.00	200.00 \pm 0.00
	7	188.60 \pm 29.61	200.00 \pm 0.00	199.94 \pm 1.14	200.00 \pm 0.00
	10	177.19 \pm 52.83	199.75 \pm 3.50	200.00 \pm 0.00	200.00 \pm 0.00
Acrobot	1	-130.60 \pm 55.08	-133.14 \pm 60.80	-81.35 \pm 22.40	-77.26 \pm 18.03
	4	-93.20 \pm 32.58	-94.21 \pm 47.20	-94.80 \pm 46.08	-83.12 \pm 23.31
	7	-96.92 \pm 34.51	-95.08 \pm 46.67	-95.75 \pm 46.57	-82.56 \pm 20.95
	10	-95.09 \pm 33.33	-77.22 \pm 18.51	-94.32 \pm 46.51	-78.91 \pm 15.76
Mountain Car	1	-136.76 \pm 34.44	-100.97 \pm 12.54	-115.48 \pm 36.35	-101.55 \pm 10.32
	4	-133.25 \pm 29.97	-99.29 \pm 8.33	-143.58 \pm 50.08	-101.35 \pm 10.63
	7	-127.34 \pm 29.15	-100.65 \pm 9.36	-128.96 \pm 46.13	-99.90 \pm 7.97
	10	-123.14 \pm 28.26	-100.48 \pm 8.14	-120.05 \pm 36.66	-100.83 \pm 11.40
HalfCheetah	4	-493.62 \pm 246.58	734.01 \pm 84.59	1008.14 \pm 280.42	4515.70 \pm 549.49
	11	637.57 \pm 1708.10	-375.22 \pm 291.13	226.06 \pm 307.87	4280.65 \pm 1119.93
	18	2705.01 \pm 2273.00	343.58 \pm 159.66	1084.26 \pm 317.02	4749.43 \pm 149.04
	25	3718.58 \pm 1856.22	502.29 \pm 375.78	869.55 \pm 447.90	4840.07 \pm 95.36
Hopper	4	50.57 \pm 0.95	3571.98 \pm 6.35	3065.21 \pm 147.79	3614.22 \pm 7.17
	11	1025.84 \pm 266.86	3572.30 \pm 12.03	3502.71 \pm 14.54	3615.00 \pm 4.32
	18	1949.09 \pm 500.61	3230.68 \pm 4.58	3201.05 \pm 6.74	3600.70 \pm 4.24
	25	3383.96 \pm 657.61	3331.05 \pm 3.55	3458.82 \pm 5.40	3560.85 \pm 3.09
Walker	4	32.18 \pm 1.25	3648.17 \pm 327.41	4945.90 \pm 65.97	4877.98 \pm 2848.37
	11	5946.81 \pm 1733.73	4723.44 \pm 117.18	6139.29 \pm 91.48	6850.27 \pm 39.19
	18	1263.82 \pm 1347.74	4184.34 \pm 485.54	5288.68 \pm 37.29	6964.68 \pm 46.30
	25	1599.36 \pm 1456.59	4368.15 \pm 267.17	4687.80 \pm 186.22	6832.01 \pm 254.64
Ant	4	1611.75 \pm 359.54	-2052.51 \pm 49.41	-5743.81 \pm 723.48	3186.80 \pm 903.57
	11	3065.59 \pm 635.19	-4462.70 \pm 53.84	-6252.19 \pm 409.42	3306.67 \pm 988.39
	18	2597.22 \pm 1366.57	-5148.62 \pm 37.80	-3067.07 \pm 177.20	3033.87 \pm 1460.96
	25	3235.73 \pm 1186.38	-5122.12 \pm 703.19	-3271.37 \pm 226.66	4132.90 \pm 878.67
Humanoid	80	1397.06 \pm 1057.84	5093.12 \pm 583.11	5096.43 \pm 24.96	10200.73 \pm 1324.47
	160	3655.14 \pm 3714.28	5120.52 \pm 17.07	5412.47 \pm 19.53	10119.80 \pm 1254.73
	240	5660.53 \pm 3600.70	5192.34 \pm 24.59	5145.94 \pm 21.13	10361.94 \pm 61.28
Task	Dataset size	Behavioral cloning	Ours ($\lambda = 0$)	Ours ($\lambda = 10^{-3}$)	Ours ($\lambda = 10^{-2}$)
Reacher	4	-10.97 \pm 7.07	-67.23 \pm 88.99	-32.37 \pm 39.81	-46.72 \pm 82.88
	11	-6.23 \pm 3.29	-6.06 \pm 5.36	-6.61 \pm 5.11	-9.26 \pm 21.88
	18	-4.76 \pm 2.31	-8.25 \pm 21.99	-5.66 \pm 3.15	-5.04 \pm 2.22

- **Paper Link:** <https://arxiv.org/pdf/1606.03476.pdf> - **2016 NIPS paper**
- **Useful blog links:**
 - <https://medium.com/@sanketgujar95/generative-adversarial-imitation-learning-266f45634e60>
 - https://hollygrimm.com/rl_gail
- **Code links:**
 - <https://github.com/hollygrimm/gail-mujoco>
 - <https://github.com/andrewliao11/gail-tf>
- **Youtube links:**
 - **Imitation learning tutorial:**
<https://www.youtube.com/watch?v=6rZTaboSY4k>
 - **Author presentation:**
<https://www.youtube.com/watch?v=bcnCo9RxhB8>